

Intro to Tantivy

A full-text search engine library written in Rust





Hello 🙌

- Adrien Guillo (@guilload)
- Software engineer
- Co-founder at Quickwit
- Two-time Recurser (Fall '14, Summer 1'17)



What is Tantivy?

Dictionary

Definitions from [Oxford Languages](#) · [Learn more](#)



tan·tiv·y

/tan' tivē/

ARCHAIC

noun

a rapid gallop or ride.

exclamation

used as a hunting cry.



What is Tantivy?


- Full-text search library inspired by **Lucene**
- Fully-featured (phrase queries, faceting, ranking, aggregations, ...)
- Mature and production-ready
- Open-source (MIT)

Tantivy is fast!



Search Benchmark - Average query latency in μ s

■ tantivy 0.16 ■ pisa 0.8.2 ■ lucene 8.10.1 ■ bleve 0.8

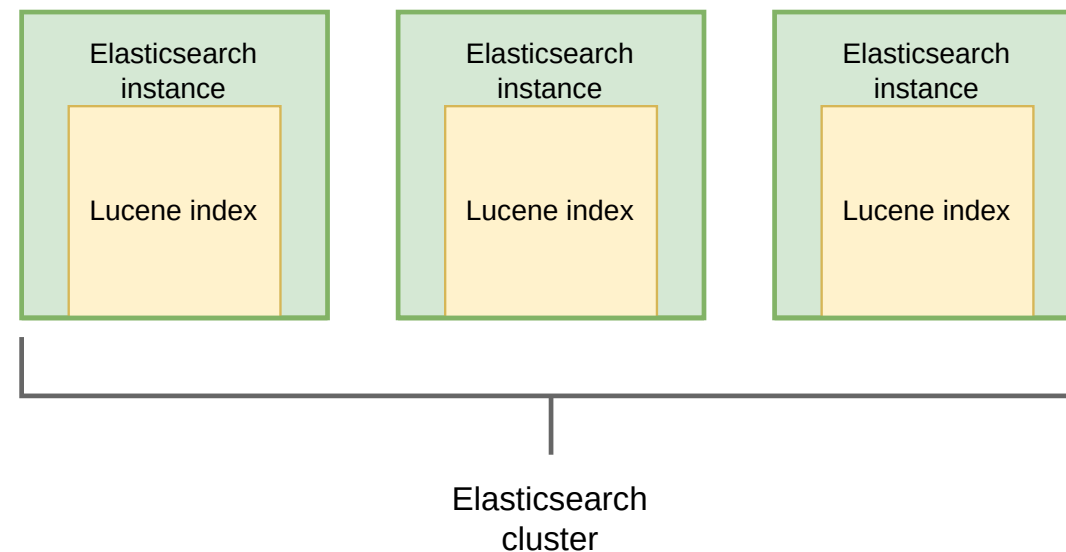
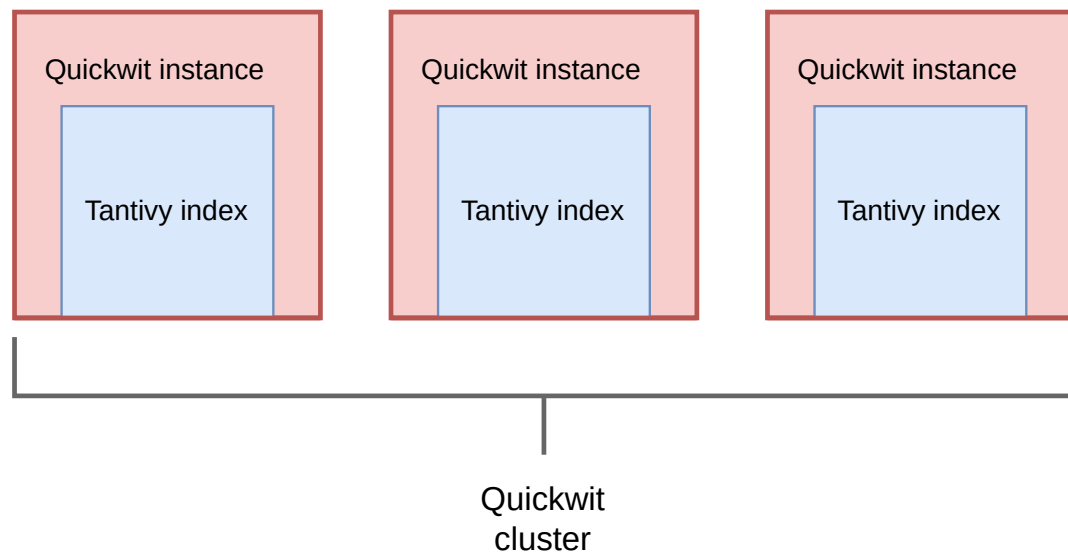


Benchmark source

“Merely” a library



- Provides the building blocks for implementing search functionalities
- NOT a distributed search engine



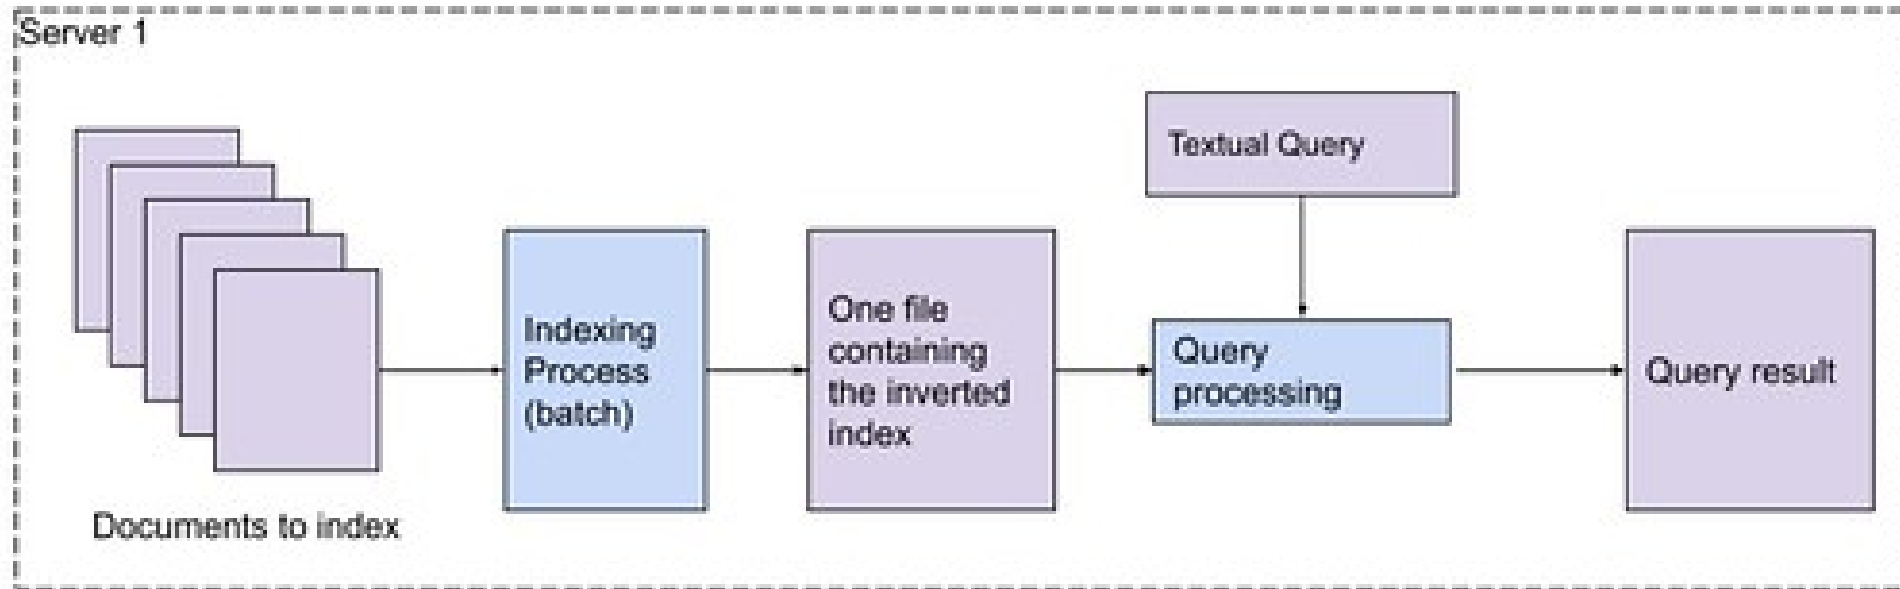


How does Tantivy work?





Search architecture





Anatomy of a search index

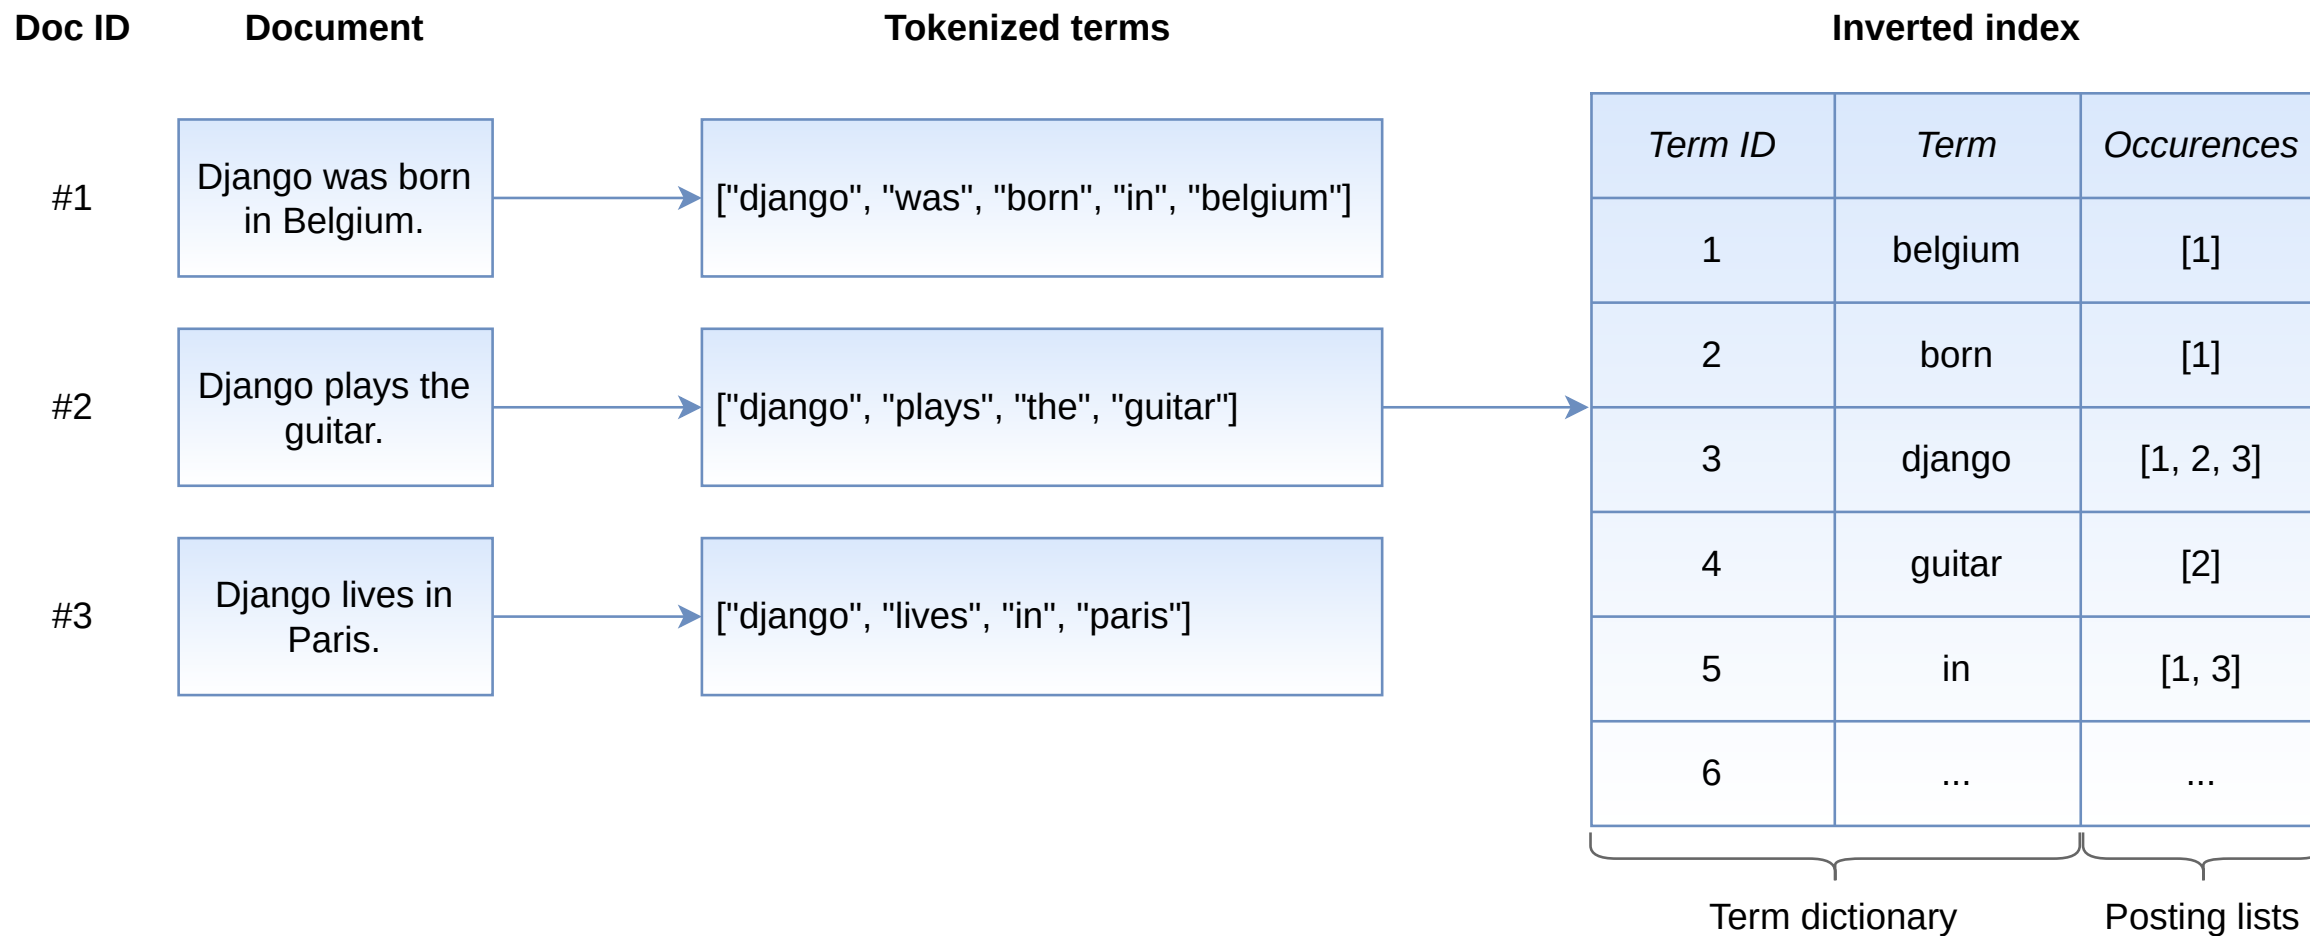
What are the data structures that compose a search index?



Anatomy of a search index

- Inverted index

Inverted index





Anatomy of a search index

- Inverted index
 - Posting lists



Posting list


- **Sorted** list of document IDs
- Efficient data structure for processing boolean queries

"django" → [1, 2, 3]



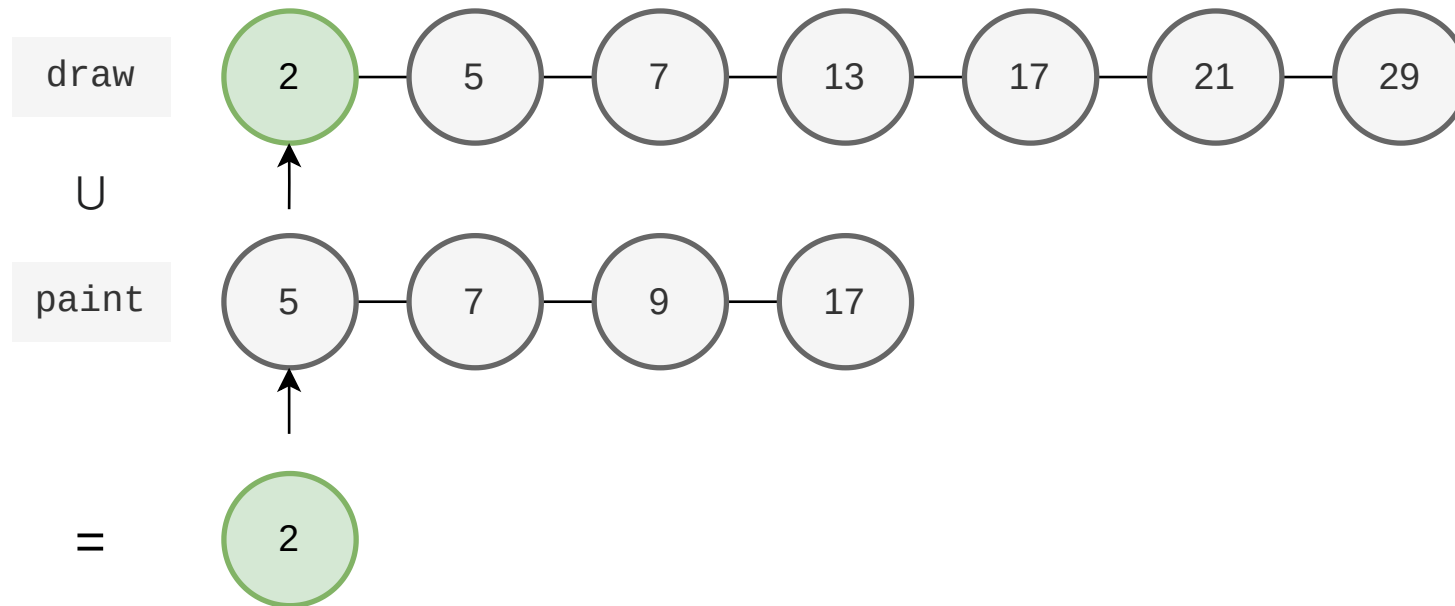
OR boolean query

- draw OR paint
- union of posting lists



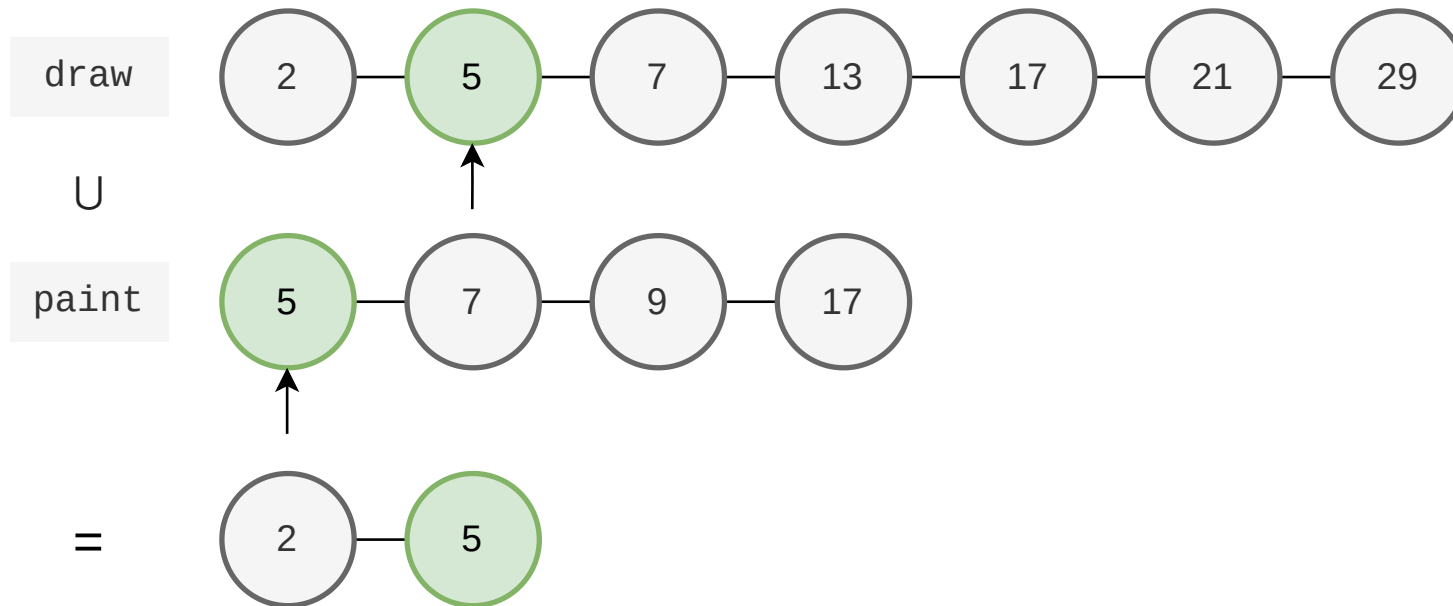


OR boolean query



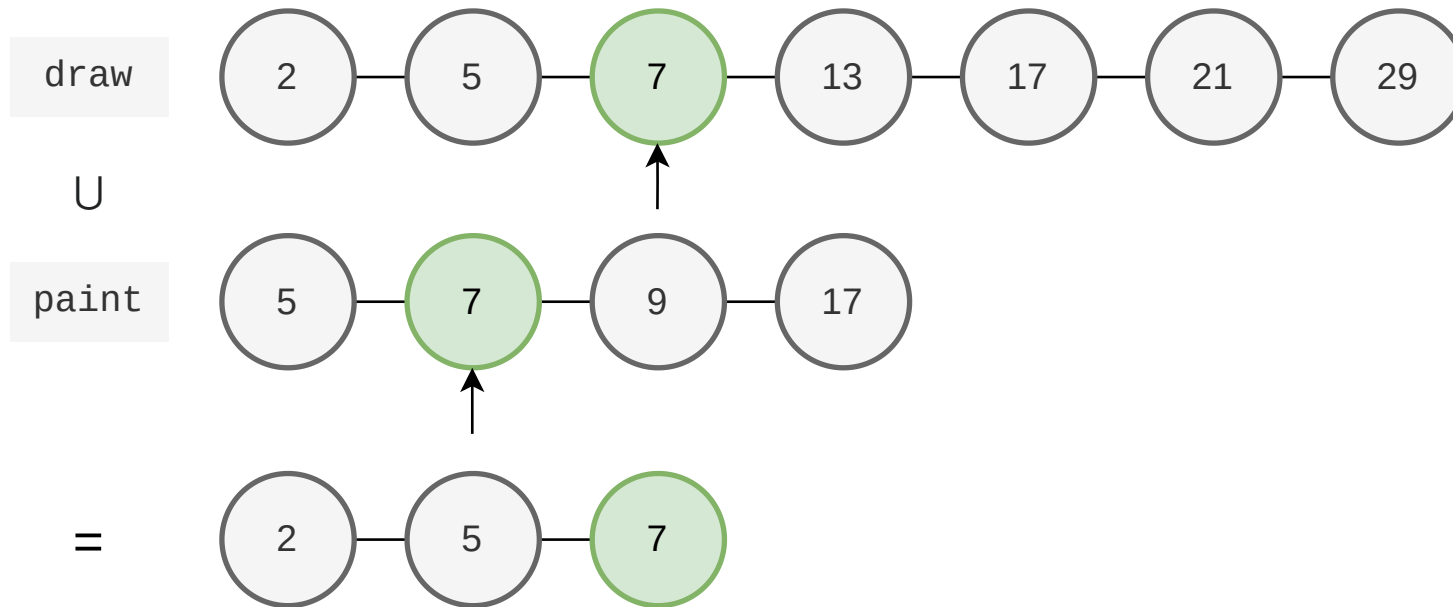


OR boolean query



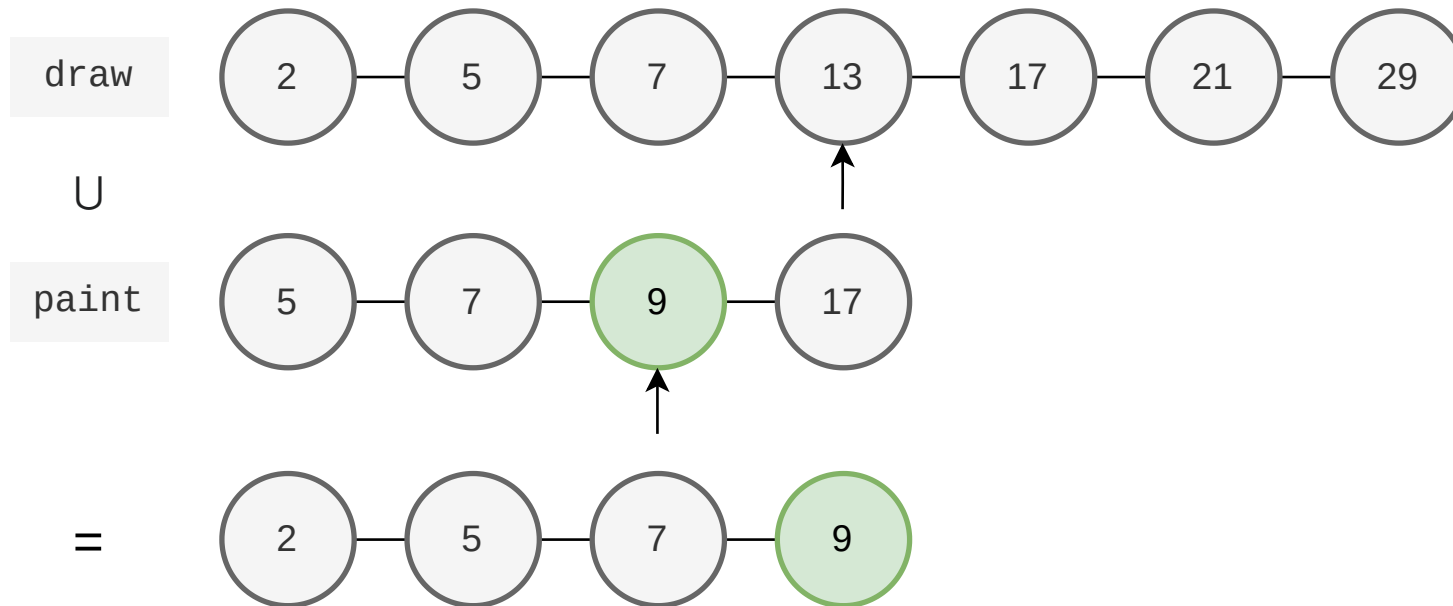


OR boolean query



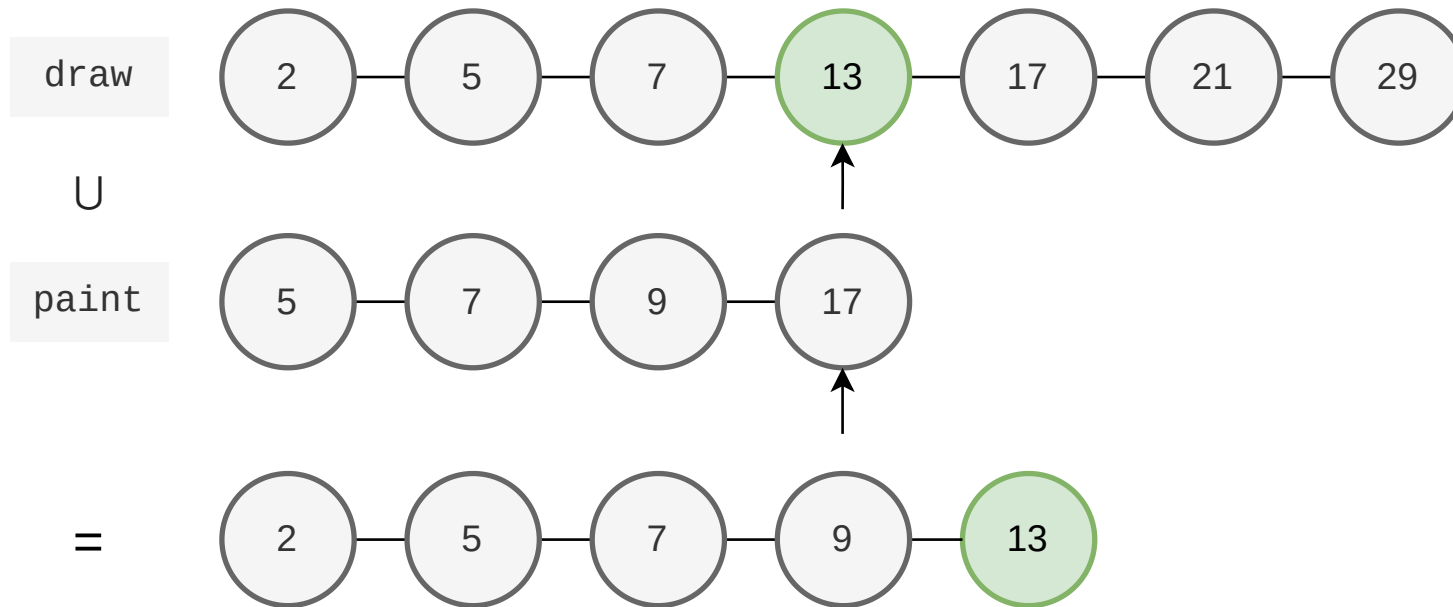


OR boolean query



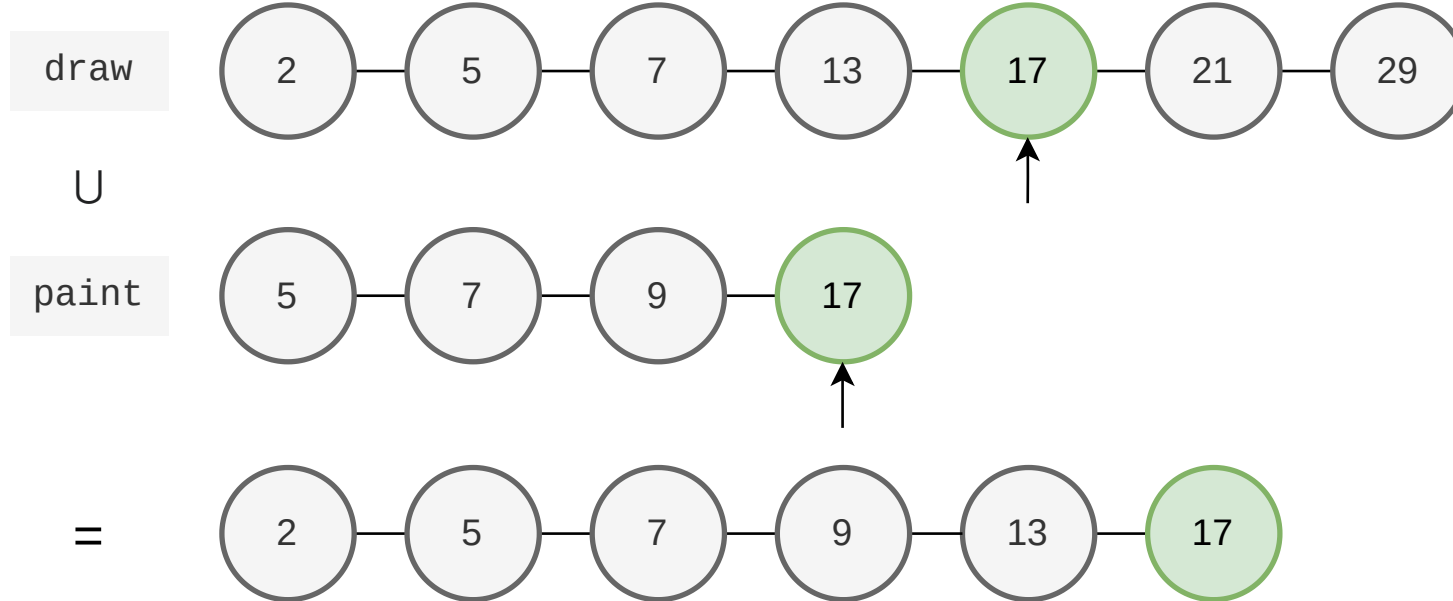


OR boolean query



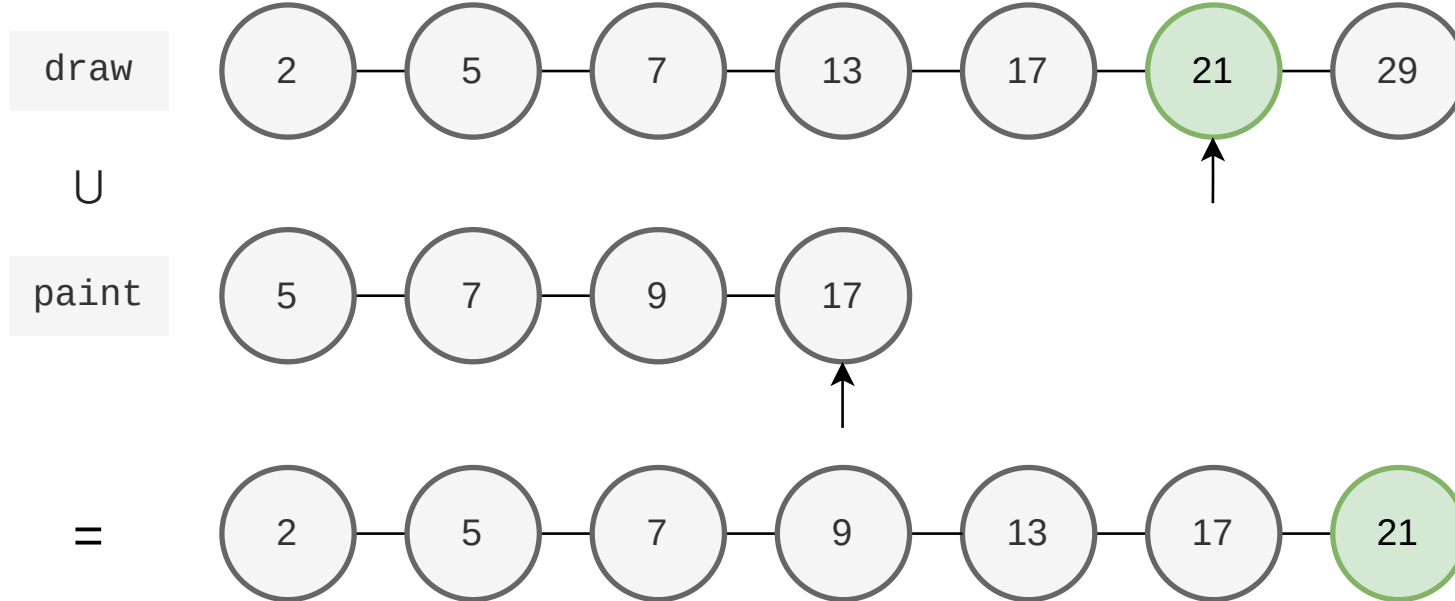


OR boolean query



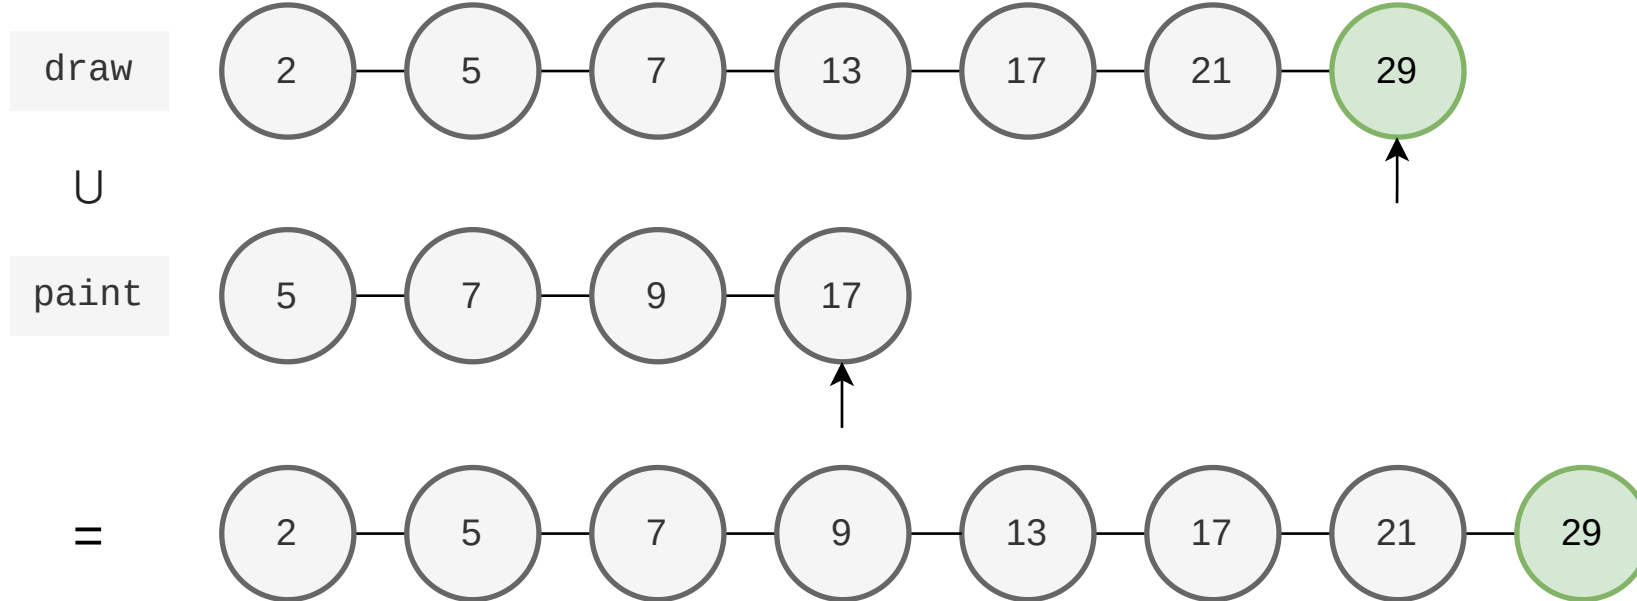


OR boolean query





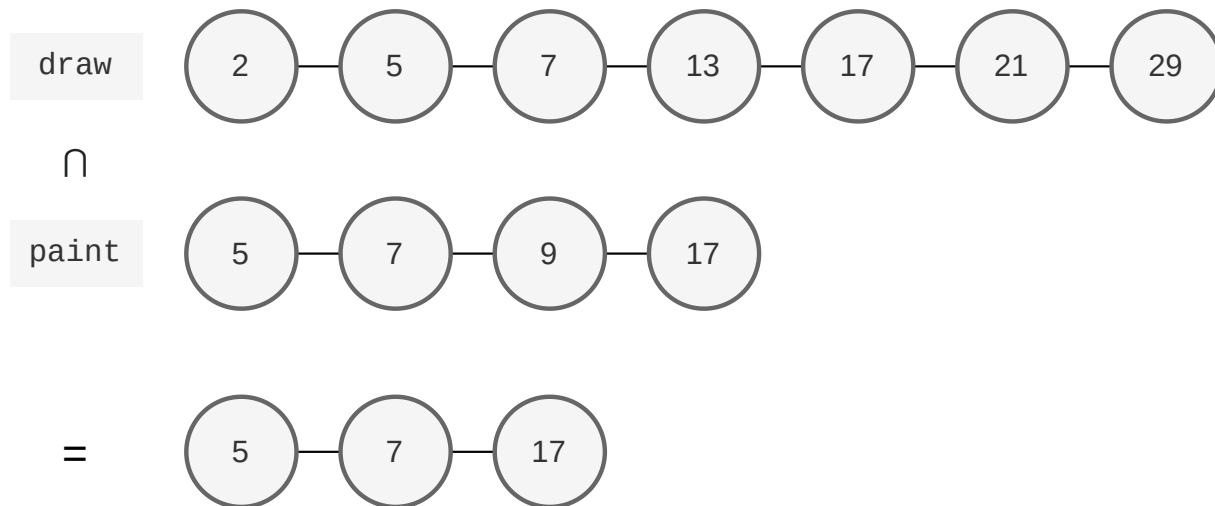
OR boolean query





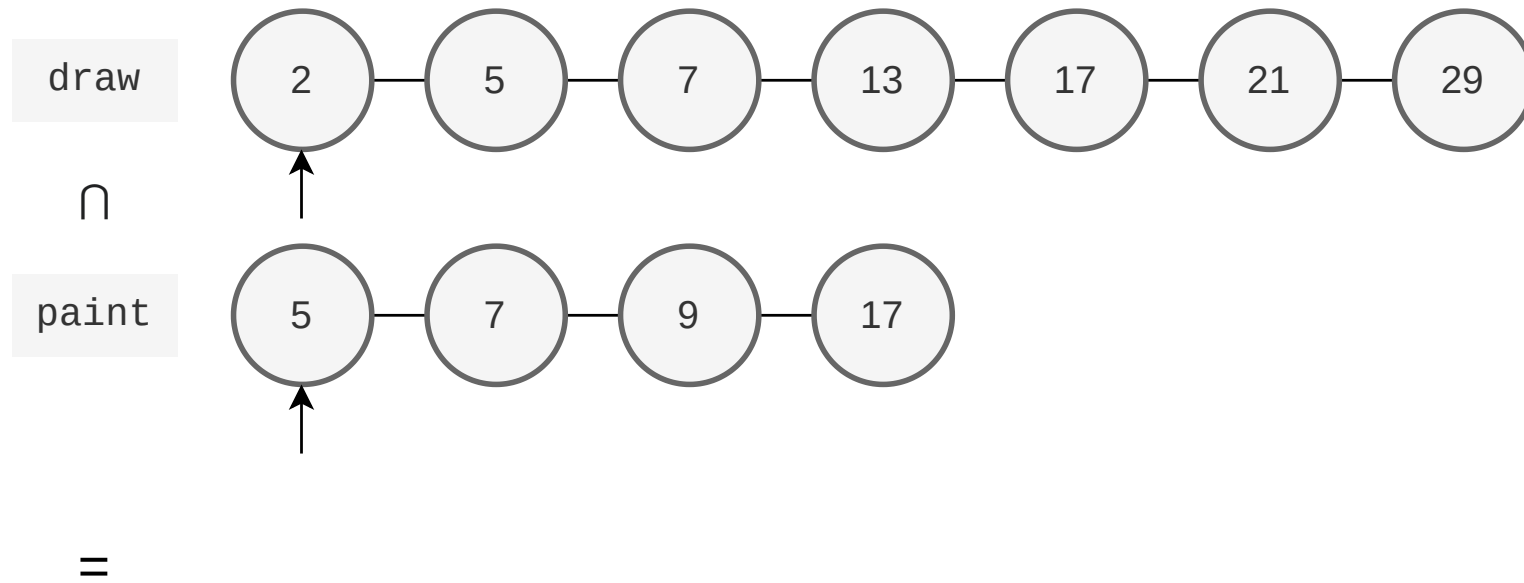
AND boolean query

- draw AND paint
- intersection of posting lists



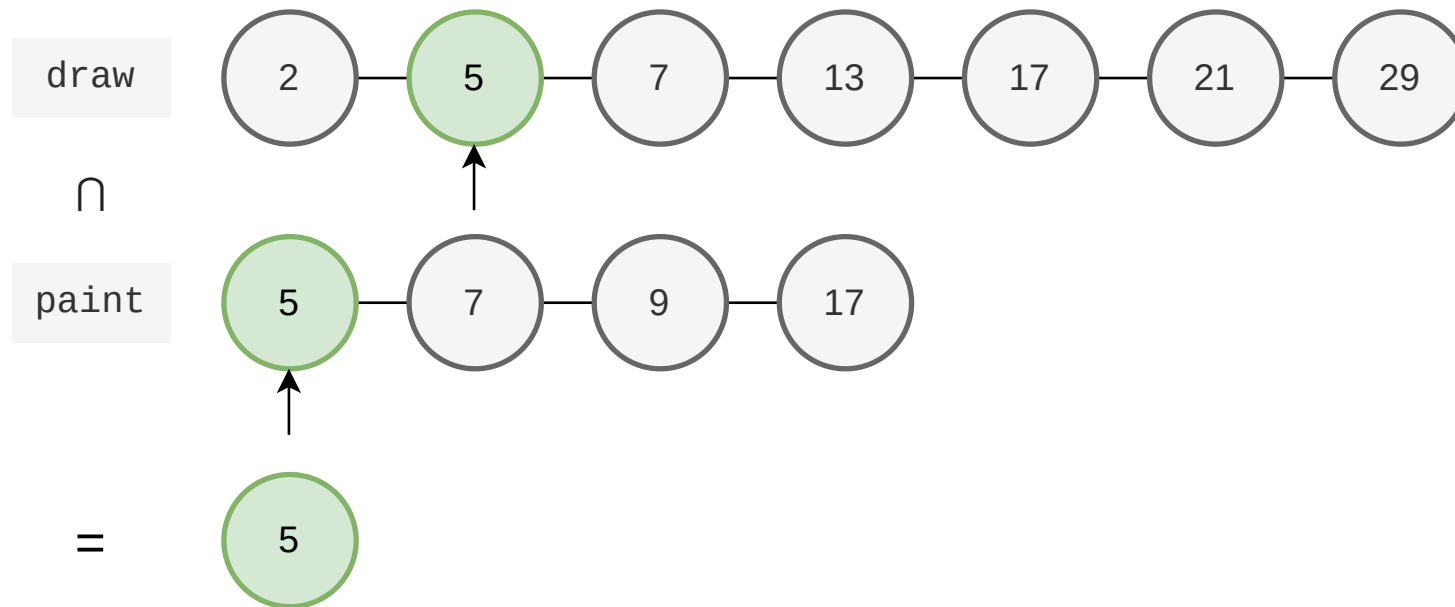


AND boolean query



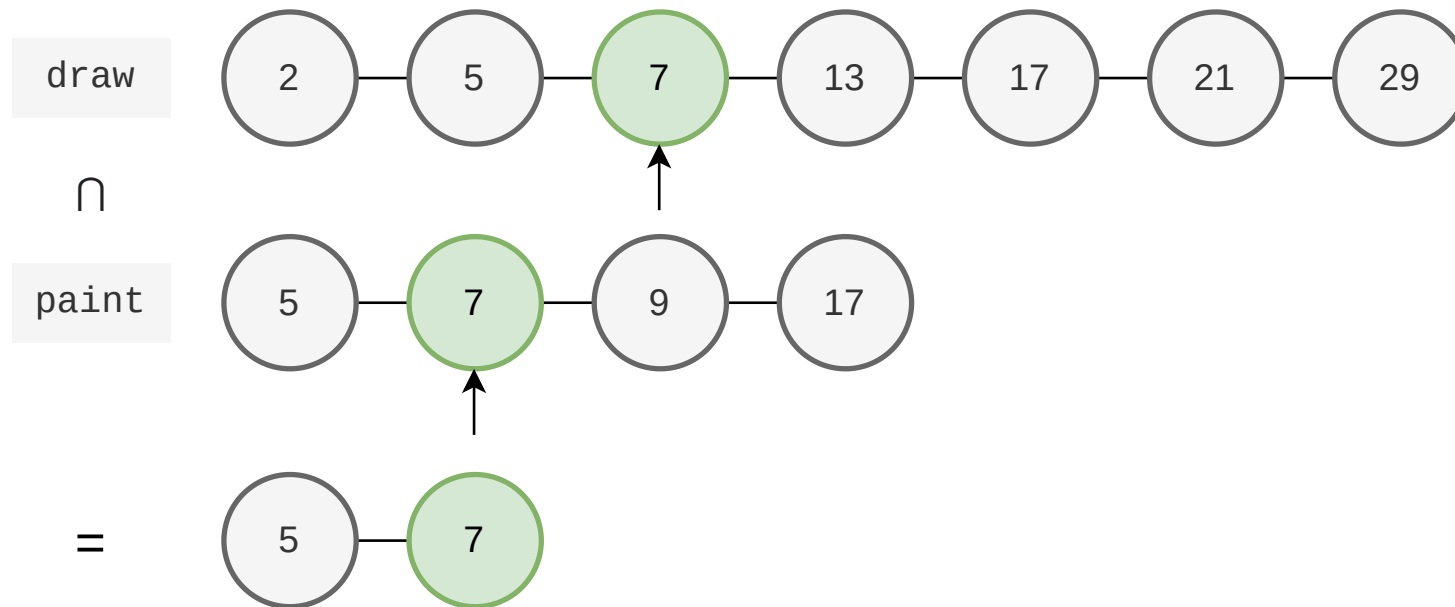


AND boolean query






AND boolean query



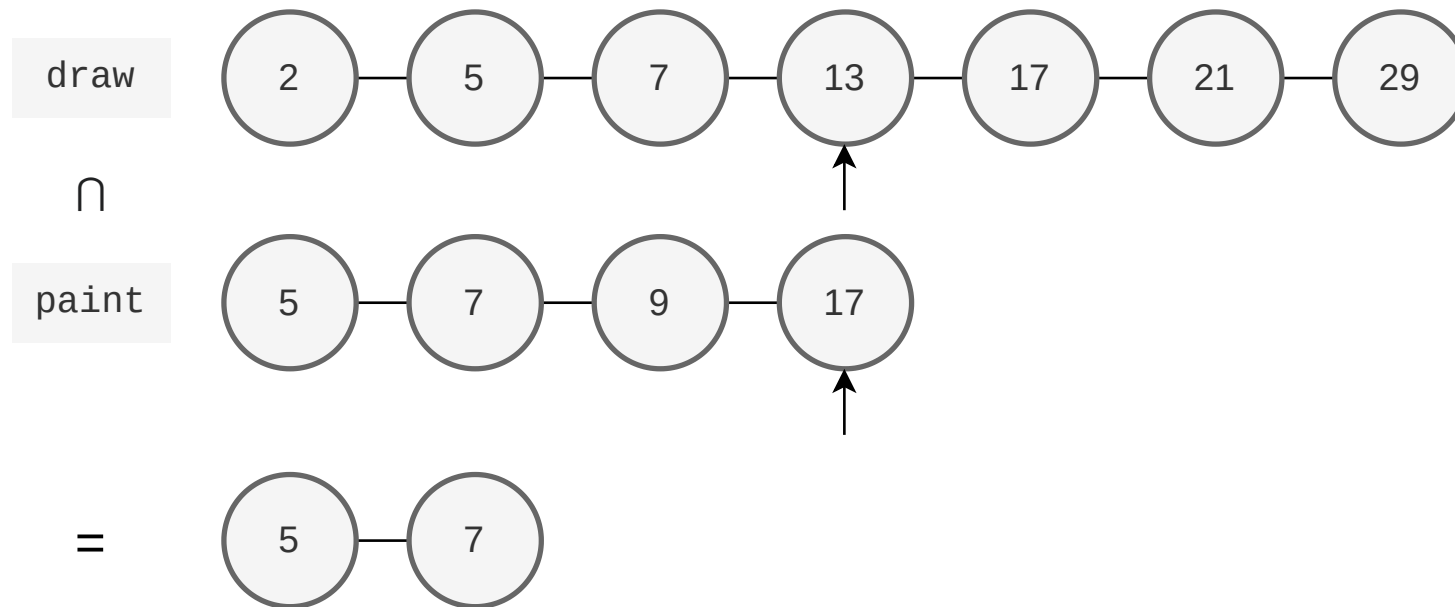


AND boolean query



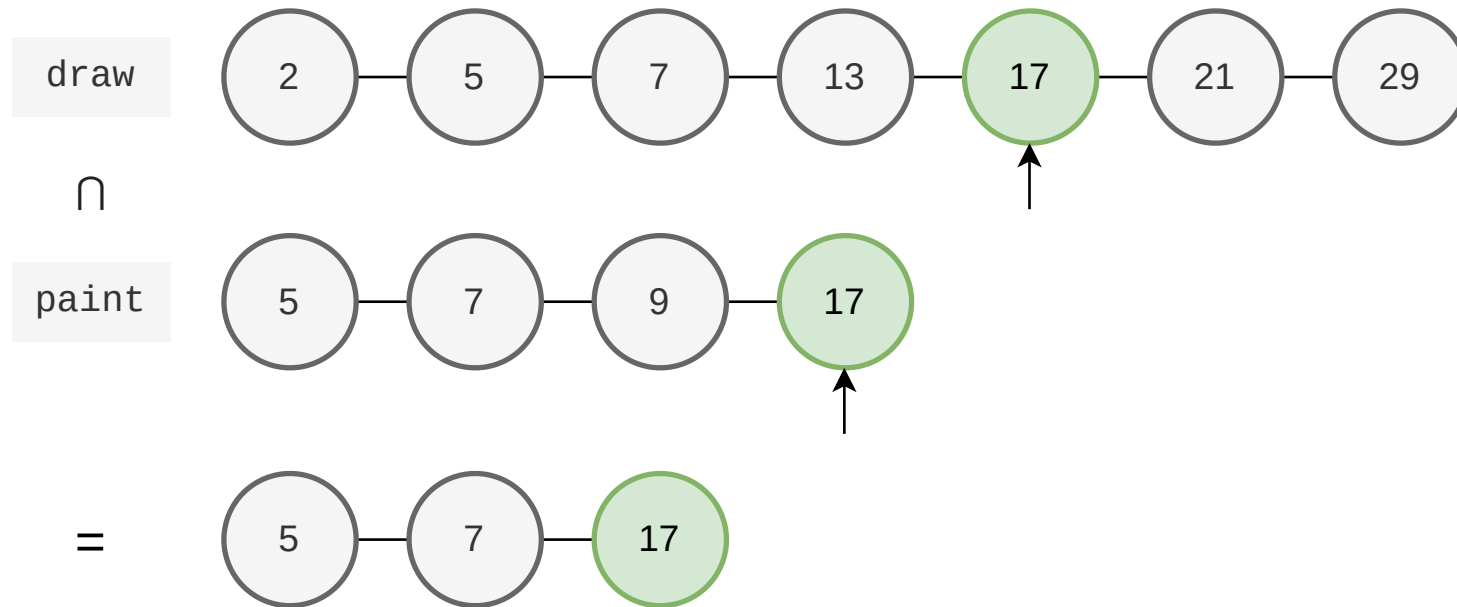


AND boolean query





AND boolean query





Posting lists compression





Delta encoding

- Document IDs are large numbers
- but differences between subsequent document IDs in posting lists can be small
- especially true for frequent terms

Example:

$[12, 17, 21, 36, , \dots] \rightarrow [12, (17 - 12), (21 - 17), (36 - 21), \dots] \rightarrow [12, 5, 4, 15, \dots]$



Bit packing

- Bit packing reduces the number of bits necessary to serialize a sequence of integers
- Performs better than dictionary-based compression schemes



Bit packing algorithm

1. Group integers into blocks of constant size
2. Find b such that all integers in the block are strictly smaller than 2^b
3. Concatenate b least significant bits of each integer



Bit packing example 1

Number	Binary representation ($b = 3$)
4	100
7	111
3	011
2	010

4 integers \rightarrow 12 bits



Bit packing example 2

Number	Binary representation (b = 4)
4	100
9	1001
3	011
2	010

4 integers → 16 bits



Bit packing at the speed of light

- SIMD instructions
- Operate on 4 integers at a time
- 4 billion integers/second!



Anatomy of a search index

- Inverted index
 - Posting lists
 - Term dictionary



Term dictionary

- A mapping from term to an offset in the postings file
 - `&[u8]` → `TermInfo`

```
//  
// Code sample extracted from `tantivy/src/postings/term_info.rs`.  
//  
pub struct TermInfo {  
    /// Byte range of the posting list within the postings (`.idx`) file.  
    pub postings_range: Range<usize>,  
}
```




Term dictionary implementation

- Hashmap family?
 - fast lookups 😊
 - one or two random I/O accesses 😊
- Trie family?
 - more expensive lookups 😞
 - more random I/O accesses 😞
 - sorted range of keys 😊
 - compress well 😊
 - intersection with automaton (Levenshtein, regex) 😊




Trie





Finite state transducer (FST)

- More compact than a trie
- More CPU intensive
- Supports same operations as a trie
- Based on the `fst` crate from [@BurntSushi](#) (author of ripgrep)





Anatomy of an invertex index

- Inverted index
 - Posting lists
 - Term dictionary
 - Sufficient?



Phrase search

	washington state	"washington state"
<i>Washington is a US state.</i>	✓	
<i>Washington State University is in Pullman.</i>	✓	✓



Positional indexing

- Record the position of each term within the document
- Enable phrase and proximity search
- Increase substantially the size of the index



Scoring

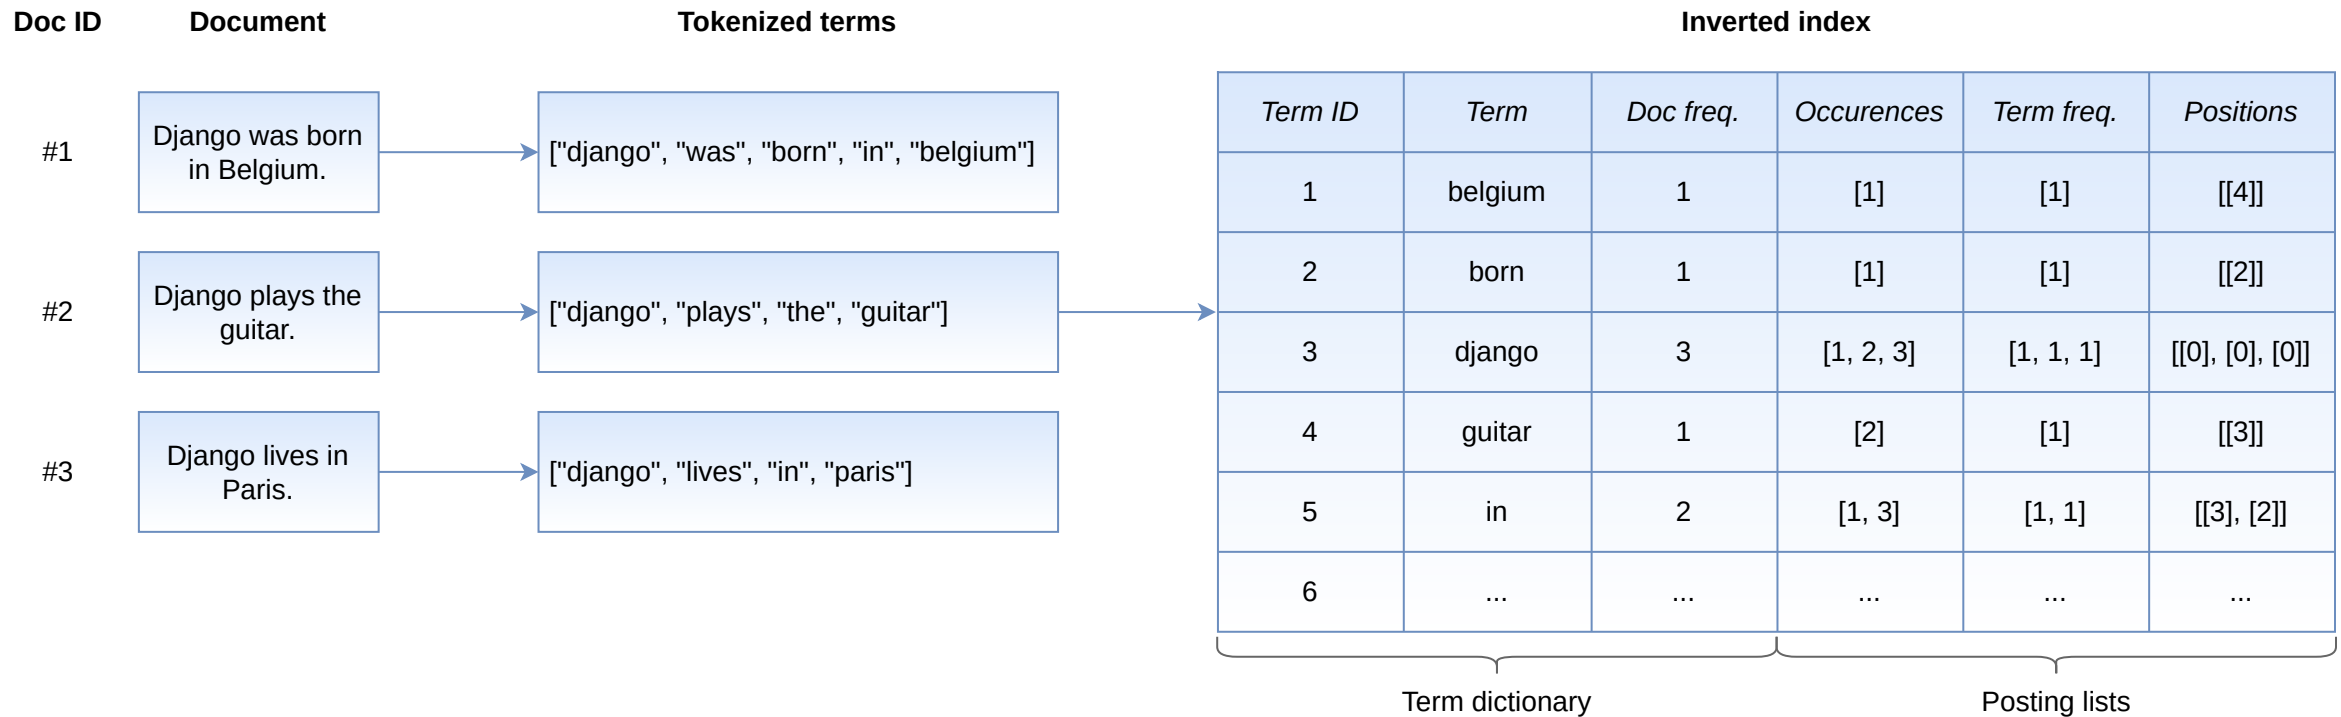
- Probabilistic scoring functions (tf-idf, BM25)



Probabilistic scoring functions

- Term frequency
 - # occurrences of term in document
 - stored in posting lists
- Document frequency
 - # documents where term occurs
 - stored in term dictionary

Inverted index with frequencies and positions





Dictionary with doc frequencies and positions

```
//  
// Code sample extracted from `tantivy/src/postings/term_info.rs`.  
//  
pub struct TermInfo {  
    /// Number of documents in the segment containing the term.  
    pub doc_freq: u32,  
  
    /// Byte range of the posting list within the postings (`.idx`) file.  
    pub postings_range: Range<usize>,  
  
    /// Byte range of the positions of this term within the positions (`.pos`) file.  
    pub positions_range: Range<usize>,  
}
```




Incremental indexing

- Inverted index is a very a compact data structure not designed for in-place updates



Segments

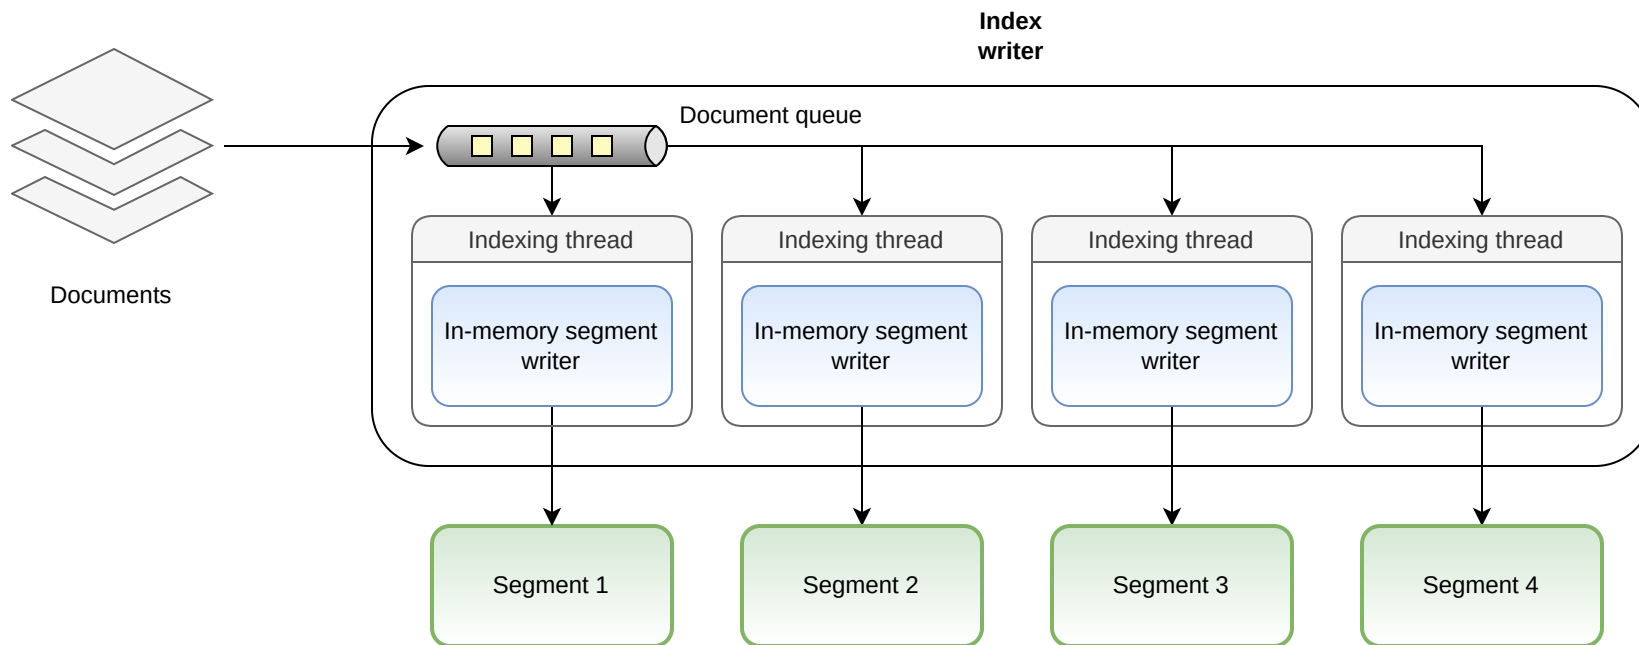
- User allocates a RAM budget (e.g. 500MB)
- Append new documents to in-memory buffer
- When the budget is reached, serialize documents to disk in the index format

→ This **small, immutable, and independent** index is called a **segment**.

→ An index is actually a **collection of segments**.



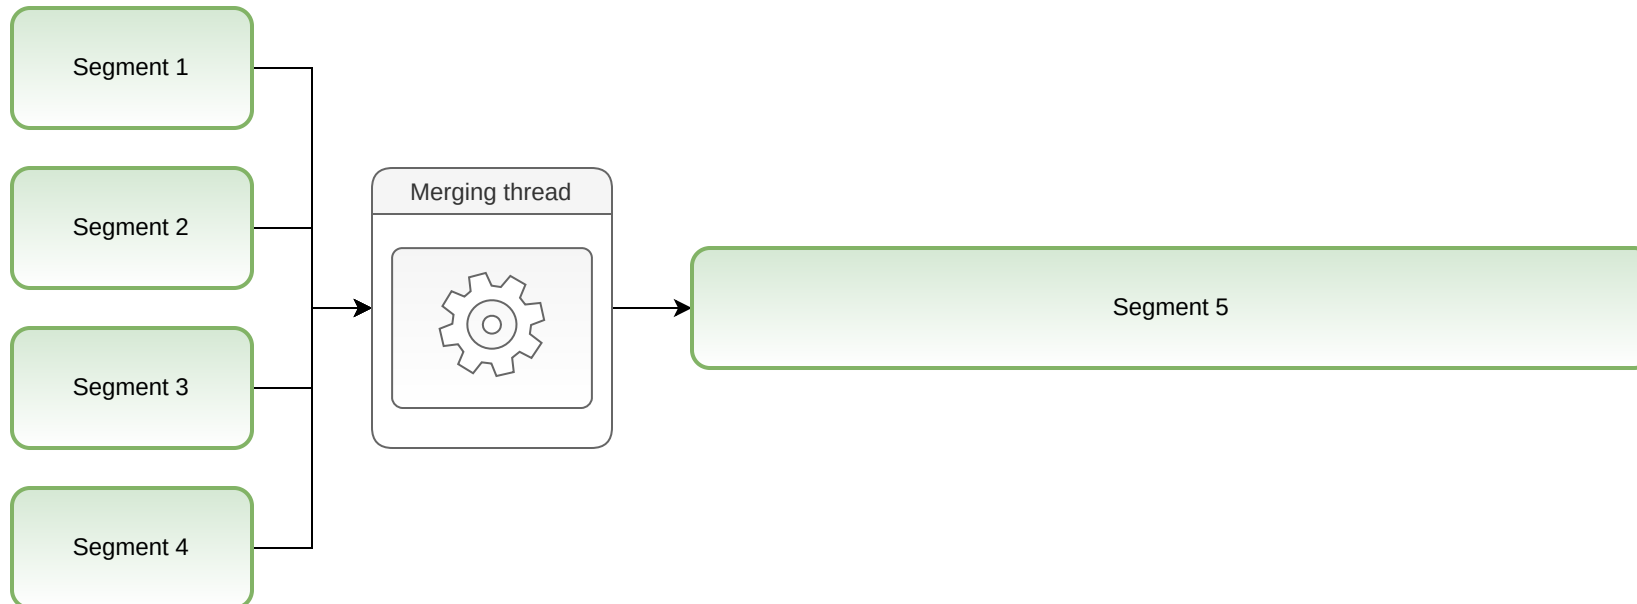
Multithreaded index writer



Segment merges



- Merge multiple segments into a larger one
- Costly operation
- Trade-off between time to search and write amplification





Let's index Wikipedia!





Indexing Wikipedia

- 5M+ Wikipedia articles (JSON)
- 8.9 GB (uncompressed)

```
{  
  "url": "https://en.wikipedia.org/wiki?curid=48689927",  
  "title": "The Princess and the Clown",  
  "body": "The Princess and the Clown (French: La princesse aux clowns) is a 1924 French silent film directed by André Hugon and starring Huguette Duflos, Charles de Rochefort and Magda Roche."  
}
```



Indexing Wikipedia

1. `create_index()`

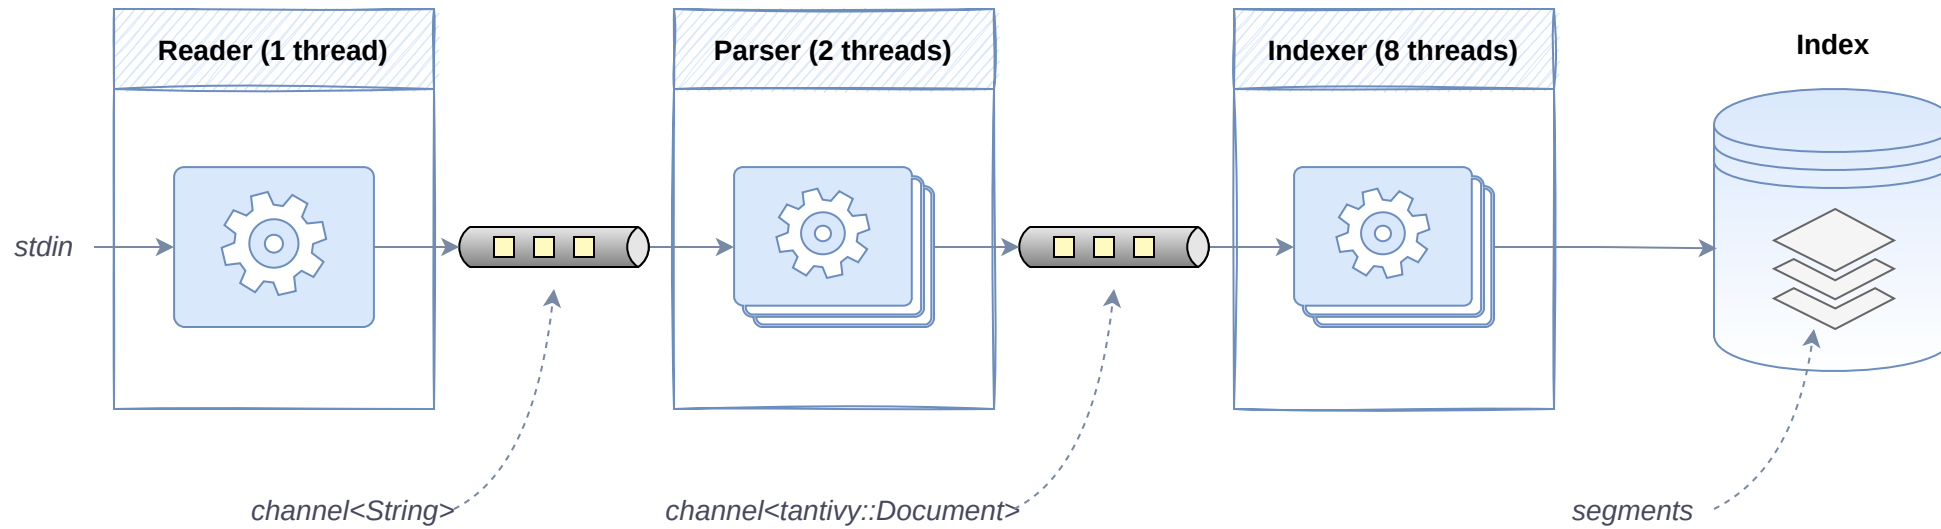


Indexing Wikipedia

1. `create_index()`
2. `append_docs()`



Indexing pipeline





Indexing Wikipedia

1. `create_index()`
2. `append_docs()`
3. `search_index()`



Indexing Wikipedia

1. `create_index()`
2. `append_docs()`
3. `search_index()`
4. `merge_segments()`



Indexing Wikipedia

1. `create_index()`
2. `append_docs()`
3. `search_index()`
4. `merge_segments()`
5. `run_gc()`



Common gotchas ⚠️

- Committing too often
- Acknowledging writes before committing
- Calling sync APIs in async runtime



Join the community!

- [GitHub](#)
- [Discord](#)



Questions?



Feedback?

adrien@quickwit.io